

Intra-daily Volume Modeling and Prediction for Algorithmic Trading

Christian T. Brownlees* Fabrizio Cipollini[†] Giampiero M. Gallo[‡]

Current Version: February 2010

Abstract

The explosion of algorithmic trading has been one of the most prominent recent trends in the financial industry. Algorithmic trading consists of automated trading strategies that attempt to minimize transaction costs by optimally placing orders. The key ingredient of many of these strategies are intra-daily volume proportions forecasts. This work proposes a dynamic model for intra-daily volumes that captures salient features of the series such as time series dependence, intra-daily periodicity and volume asymmetry. Moreover, we introduce a loss functions for the evaluation of proportions forecasts which retains both an operational and information theoretic interpretation. An empirical application on a set of widely traded index Exchange Traded Funds shows that the proposed methodology is able to significantly outperform common forecasting methods and delivers more precise predictions for Volume Weighted Average Price trading.

*Department of Finance, Stern School of Business, NYU and Dipartimento di Statistica “G. Parenti”, Università di Firenze, Italy. e-mail: ctb@stern.nyu.edu

[†]Dipartimento di Statistica “G. Parenti”, Università di Firenze, Italy. e-mail: cipollini@ds.unifi.it

[‡]Corresponding author: Dipartimento di Statistica “G. Parenti”, Università di Firenze, Viale G.B. Morgagni, 59 - 50134 Firenze – Italy. e-mail: gallog@ds.unifi.it

First draft: Dec. 2008. We would like to thank Torben Andersen and two anonymous referees for their comments and concerns. Thanks are due also Robert Almgren, Bruno Biais, Claude Courbois, Rob Engle, Thierry Foucault, Terry Hendershott, Eric Ghysels, Farhang Farazmand, Dennis Kristensen, Albert Menkveld, Giovanni Urga, seminar participants in the Columbia Econometrics Workshop and at the Goldman Sachs GSET Strategists Weekly seminar; conference participants in the Chicago London Conference, Dec. 5–6, 2008; the First FBF–IDEI-R Conference on Investment Banking and Financial Markets, Toulouse, March 26–27, 2009; The SOFIE 1st European Conference Geneva, 16–18 June, 2009; and (EC)²–Conference Aarhus, 18–19 December, 2009. The usual disclaimer applies.

1 Introduction

The last few years have witnessed a widespread development of automated order execution systems, typically known in the financial industry as *algorithmic* (or *algo*) *trading*. Such algorithms aim at enhancing order execution by strategically submitting orders: computer-based pattern recognition allows for instantaneous information processing and for subsequent action taken with limited (if any) human judgement and intervention. The impact of such an innovation has been quite substantial: Chordia *et al.* (2008) recognize that algo trading is the main determinant of the increase in volume and the reduction in the average trade size observed over the recent years. In October 2006, the NYSE has boosted a mixed system of electronic and face-to-face auctions which brings automated trades to more than 70% of total trades according to some 2009 estimates.

One of the industry's concerns is to split orders in order to seek better price execution, while retaining transparency of the procedures *vis à vis* the client. The *Volume Weighted Average Price* (VWAP) (cf. Madhavan (2002)) of an asset is a well established benchmark to fulfill this goal: it is calculated at the end of the trading day as an average of intra-daily transaction prices weighted by the corresponding share of traded volume to total daily volume. In the simplest of the scenarios, customers can be guaranteed that their order will be executed at the VWAP and the algorithm can be tuned to spread that order during the day with the goal of achieving an average execution price close to the VWAP. Such order execution procedure is fairly popular and is commonly named *VWAP trading*. From an econometric point of view, it is interesting to note that the strategy needs to be based on accurate predictions of intra-daily volume proportions.

This paper deals with modeling and forecasting intra-daily volume time series with an application to VWAP trading. Despite the widespread use of ultra-high frequency data in econometrics for more than 15 years, there has been little investigation of the features of intra-daily volumes and there is no well established econometric methodology for fore-

casting them. The main contribution of this paper is to turn a descriptive analysis of the patterns of intra-daily volume time series into the design of a novel dynamic specification which replicates the empirical regularities. In particular, daily average dynamics are strongly persistent, whereas intra-daily patterns are characterized by clustering around a stable intra-daily U-shaped periodic component. Joined with the fact that volumes are non-negative quantities, such evidence suggests a *Component Multiplicative Error Model* (CMEM), an extension of the Multiplicative Error Model (MEM, Engle (2002)). We do not choose a specific distribution for the error term, preferring to work with a semiparametric specification, and we estimate all parameters jointly by Generalized Method of Moments. We evaluate volume forecast performance in the perspective of algorithmic trading proposing a loss function, the *Slicing Loss* function, suitable for proportions.

The empirical analysis is carried out on modeling volume turnover of three liquid Exchange Traded Funds (ETF) which replicate the behavior of major US stock indices: SPDR S&P 500 (SPY – S&P 500), Diamonds (DIA – Dow Jones) and PowerShares QQQ (QQQQ – NASDAQ) between 2001 and 2006. We use different metrics to measure precision gains in forecasting volumes, volume proportions and VWAP. Results show that our model significantly outperforms a simple benchmark (common among practitioners), signaling that order execution can be enhanced by an econometric model molded after stylized facts.

Our paper fits in the large literature on intra-daily trading activity modeling. In particular, our model extends the logic of the component GARCH model for intra-daily volatility suggested by Engle *et al.* (2007) and has some connections with P-GARCH models introduced by Bollerslev and Ghysels (1996). Białkowski *et al.* (2008) concentrate on volume dynamics of several stocks, isolating common and idiosyncratic components. Gouriéroux *et al.* (1999) fix a volume level and model the time needed to trade it. Andersen (1996) models volumes together with volatility as a function of an underlying latent information flow. The paper is also related to a recent strand of literature which shows how economet-

ric methodology can be applied to the analysis and reduction of transaction costs, such as Taylor (2002), Engle *et al.* (2006), Bikker *et al.* (2008) and Härdle *et al.* (2009) among others.

The paper is organized as follows. We start from stylized facts (Section 2) to motivate the Component MEM (Section 3). Section 3.3 gives details on the estimation procedure. The empirical application is divided up between model estimation and diagnostics (Section 3.4) and volume forecasting and VWAP forecast comparisons (Section 4). Concluding remarks follow (Section 5).

2 The Empirical Regularities of Intra-daily Volumes

We chose to analyze Exchange Traded Funds (ETFs), innovative financial products which allow straightforward trading in market averages as if they were stocks, while avoiding the possible idiosyncrasies of single shares. In the present framework, we count on a dataset consisting of regularly spaced intra-daily volumes and transaction price data for three popular equity index ETFs: SPY (S&P 500 ETF), DIA (Dow Jones ETF), and QQQQ (Nasdaq ETF) over a sample period between January 2002 and December 2006, excluding days with empty bins. The frequency of the intra-daily data is 15 minutes (26 intra-daily bins). Volumes are computed as the sum of all transaction volumes exchanged within each intra-daily bin: in order to correct for some possible trending behavior in volume (Andersen (1996)), we divide the original series (multiplied by 100) by the number of daily outstanding shares. The last recorded transaction price before the end of each bin is used as the reference price. The ultra-high frequency data used in the analysis are extracted from the TAQ while shares outstanding are taken from the CRSP. Details on the series handling and management are documented in Brownlees and Gallo (2006).

While similar evidence also holds for the other tickers (for which we report summary descriptive statistics only), let us focus on the empirical regularities of the SPY turnover




Figure 1: SPY Turnover Data: Original Turnover Data (top); Daily Averages (center), Intra-daily Component (bottom). January 2002 to December 2006.




Figure 2: SPY Turnover Data: Intra-daily Component (top); Intra-daily Periodic Component (center); Intra-daily Non-periodic Component (bottom). January 2002 to December 2006.




Figure 3: SPY Turnover Data: Autocorrelation Function of the Original Turnover and of the Daily Averages.

series (1247 trading days and 32422 observations) by starting with a graphic appraisal of the 15–minutes turnover (top panel of Figure 1). As with most financial time series, it clearly exhibits clustering of trading activity. Such a pattern is retained by taking daily averages (cf. second panel of Figure 1), which we interpret as evidence that the overall series clusters around a lower frequency component. Dividing each observation by the corresponding daily average we obtain the intra-daily pattern (bottom panel of Figure 1, also reproduced as the top panel of Figure 2 for ease of reference). Allowing for the possible presence of periodic and non periodic components, we compute averages by time of day (26 bins – center panel of Figure 2) which exhibit a U–shape as other intra-daily financial time series (e.g. average durations, the trading activity is higher at the opening and closing of the markets and is lower around mid-day). The ratio between these two series shows the non–periodic pattern in the bottom panel of Figure 2.

The dynamic features of these three series (daily, intra-daily periodic, and intra-daily non periodic) are summarized in the corresponding correlograms (left panel of Figure 3 for raw data; right panel for the daily averages). Focusing on the original series divided by daily averages we show marked periodicity (correlogram in the left panel of Figure 4); adjusting this series by time–of–day averages we remove periodicity but some short–lived dependence is retained (correlogram in the right panel of Figure 4).




Figure 4: SPY Turnover Data: Autocorrelation Function of the Intra-daily Component (left panel) and of the Intra-daily Non-periodic Component (right panel).

The autocorrelations of the components of the 15-minute series for all three tickers (Table 1) confirms the graphical analysis of SPY. The overall time series display relatively high levels of persistence which are also slowly decaying. The autocorrelations do not decrease by daily averaging. By dividing the overall turnover by its daily average (intra-daily component), a substantial part of dependence in the series is removed. Finally, once the intra-daily periodic component is removed, the resulting series show significant low order correlations only. Interestingly, the magnitudes of the various autocorrelations of the series are remarkably similar across the assets.

	overall		daily		intra-daily		intra-daily non-periodic	
	$\hat{\rho}_1$	$\hat{\rho}_{1 \text{ day}}$	$\hat{\rho}_{1 \text{ day}}$	$\hat{\rho}_{1 \text{ week}}$	$\hat{\rho}_1$	$\hat{\rho}_{1 \text{ day}}$	$\hat{\rho}_1$	$\hat{\rho}_{1 \text{ day}}$
SPY	0.74	0.59	0.84	0.76	0.46	0.37	0.26	0.00
DIA	0.60	0.41	0.72	0.60	0.35	0.26	0.18	0.02
QQQQ	0.69	0.53	0.77	0.66	0.53	0.45	0.29	0.01

Table 1: Autocorrelations at selected lags of the turnover time series components. The table reports the lag 1 ($\hat{\rho}_1$) and lag 26 ($\hat{\rho}_{1 \text{ day}}$) autocorrelations of the intra-daily frequency components (overall, intra-daily and intra-daily non-periodic) and lag 1 ($\hat{\rho}_{1 \text{ day}}$) and lag 5 ($\hat{\rho}_{1 \text{ week}}$) autocorrelations of the daily frequency component (daily).

3 A Multiplicative Error Model for Intra-daily Volumes

Based on the empirical regularities discussed in Section 2, we will specify the dynamics of intra-daily volumes split in one daily and two intra-daily (one periodic and one dynamic) components. Let us first establish the notation used in what follows. Days are denoted with $t \in \{1, \dots, T\}$; each day is divided into I equally spaced intervals (referred to as bins) indexed by $i \in \{1, \dots, I\}$. In order to simplify the notation we may label observations indexed by the double subscript ti with a single progressive subscript $\tau = I \times (t - 1) + i$. Correspondingly, we denote the total number of observations by N ($T \times I$).

The non-negative quantity under analysis relative to bin i of day t is denoted as x_{ti} or, alternatively, as x_τ . $\mathcal{F}_{ti-1}$ indicates the information about x_{ti} available before forecasting it. Usually, we will assume $\mathcal{F}_{t0} = \mathcal{F}_{t-1I}$ but it may be possible to include additional pieces of market opening information into $\mathcal{F}_{t0}$.

We adopt the following convention: if $\mathbf{x}_1, \dots, \mathbf{x}_K$ are (m, n) matrices then $(\mathbf{x}_1; \dots; \mathbf{x}_K)$ represents the (mK, n) matrix obtained stacking the $\mathbf{x}_t$ matrices columnwise.

3.1 Model Definition

Following the logic of Multiplicative Error Models (MEM) and in view of the motivation provided by the stylized facts in Section 2, we propose a Component MEM (CMEM)

$$x_{ti} = \eta_t \phi_i \mu_{ti} \varepsilon_{ti}.$$

The multiplicative innovation term ε_{ti} is assumed i.i.d., non-negative, with mean 1 and constant variance σ^2 :

$$\varepsilon_{ti} | \mathcal{F}_{ti-1} \sim (1, \sigma^2). \quad (1)$$

The conditional expectation of x_{ti} is the product of three multiplicative elements:

- η_t , a daily component;
- ϕ_i , an intra-daily periodic component aimed to reproduce the time-of-day pattern;
- μ_{ti} , an intra-daily dynamic (non-periodic) component.

In order to simplify the exposition, we assume a relatively simple specification for the components. If needed, the formulation proposed can be trivially generalized, for instance by including asymmetric effects, other predetermined variables, more lags and so on (see the empirical application in Section 3.4).

The daily component is structured as

$$\eta_t = \alpha_0^{(\eta)} + \beta_1^{(\eta)} \eta_{t-1} + \alpha_1^{(\eta)} x_{t-1}^{(\eta)} \quad (2)$$

where $x^{(\eta)}$ is what we name the *deflated daily volume*,

$$x_t^{(\eta)} = \frac{1}{I} \sum_{i=1}^I \frac{x_{ti}}{\phi_i \mu_{ti}}, \quad (3)$$

that is the daily average of the intra-daily volumes deflated by the intra-daily components ϕ_i and μ_{ti} .

The intra-daily dynamic component is formulated as

$$\mu_{ti} = \alpha_0^{(\mu)} + \beta_1^{(\mu)} \mu_{ti-1} + \alpha_1^{(\mu)} x_{ti-1}^{(\mu)} \quad (4)$$

where $x^{(\mu)}$ is the *deflated intra-daily volume*,

$$x_{ti}^{(\mu)} = \frac{x_{ti}}{\eta_t \phi_i}. \quad (5)$$

Furthermore, μ_{ti} is constrained to have unconditional expectation equal to 1 in order to make the model identifiable, allowing us to interpret it as a pure intra-daily dynamic component. This implies $\alpha_0^{(\mu)} = 1 - \beta_1^{(\mu)} - \alpha_1^{(\mu)}$. The starting conditions for the system we use are $\eta_0 = x_0^{(\eta)} = \frac{1}{5} \sum_{t=1}^5 \sum_{i=1}^I x_{ti}$ (that is the daily volume average over the first week of the sample) and $\mu_{10} = x_{1,0}^{(\mu)} = 1$. The intra-daily non-periodic component can be initialized at each day t with the latest quantities available, namely,

$$\mu_{t0} = \mu_{t-1I} \quad x_{t0}^{(\mu)} = x_{t-1I}^{(\mu)}.$$

In synthesis, the system nests the daily and the intra-daily dynamic components by alternating the update of the former (from η_{t-1} to η_t) and of the latter (from $\mu_{t0} = \mu_{t-1I}$ to μ_{tI}). Time-varying η_t adjusts the mean level of the series, whereas the intra-daily component $\phi_i \mu_{ti}$ captures bin-specific departures from such an average level.

Note that defining $x_{ti}^{(\mu)}$ as in (5) implies $x_{ti}^{(\mu)} = \mu_{ti} \varepsilon_{ti}$. Combining this with (1), one obtains

$$E(x_{ti}^{(\mu)} | \mathcal{F}_{t-1}) = \mu_{ti} \quad V(x_{ti}^{(\mu)} | \mathcal{F}_{t-1}) = \mu_{ti}^2 \sigma^2 \quad (6)$$

that coincide with the properties of the corresponding quantity in the usual MEM (Engle (2002)). A similar consideration can be made for $x_t^{(\eta)}$. In fact, definition (3) implies $x_t^{(\eta)} = \eta_t \bar{\varepsilon}_t$, where $\bar{\varepsilon}_t = I^{-1} \sum_{j=1}^I \varepsilon_{tj}$, and thus

$$E(x_t^{(\eta)} | \mathcal{F}_{t-1I}) = \eta_t \quad V(x_t^{(\eta)} | \mathcal{F}_{t-1I}) = \eta_t^2 \sigma^2 / I. \quad (7)$$

On this base, $x_t^{(\eta)}$ and $x_{ti}^{(\mu)}$ are adjusted versions of the observed x_{ti} 's that can be interpreted as transmitting an innovation content through the respective equations.

The intra-daily periodic component ϕ_i can be specified in various ways but here we retain

a parsimonious parameterization of ϕ_i via a Fourier (sine/cosine) representation:

$$\phi_{i+1} = \exp \left\{ \sum_{k=1}^K [\delta_{1k} \cos(fki) + \delta_{2k} \sin(fki)] \right\} \quad (8)$$

where $f = 2\pi/I$, K is the integer part of $I/2$, $\delta_{2K} = 0$ if I is even and $i = 0, \dots, I - 1$. Besides the good approximation properties of the Fourier series, what also makes this basis of functions appealing is the fact that the sine/cosine terms are orthogonal and bounded functions (as opposed, for instance, to splines) and this significantly eases the nonlinear numerical optimization required for the estimation. (also see White (2006) as well as Gallant (1981) and Gallant (1984)). Moreover, the number of terms into (8) may be considerably reduced if the periodic intra-daily pattern is sufficiently smooth, since few low frequencies harmonics may be enough. Alternatively, shrinkage type estimation may allow flexibility and parsimony of the estimated periodic component (cf. Brownlees and Gallo (2008)).

3.2 Discussion

3.2.1 Correspondence of the CMEM to the Descriptive Analysis

The daily average $\bar{x}_t = I^{-1} \sum_{i=1}^I x_{ti}$ represents a proxy of the daily component η_t . In fact, by taking its expectation conditionally on the previous day, we have

$$E(\bar{x}_t | \mathcal{F}_{t-1}) = \eta_t \frac{1}{I} \sum_{i=1}^I \phi_i E(\mu_{ti} | \mathcal{F}_{t-1}) \simeq \eta_t \frac{1}{I} \sum_{i=1}^I \phi_i = \eta_t \bar{\phi}, \quad (9)$$

where the approximate equality can be justified by noting that the non-periodic intra-daily component μ_{ti} has unit unconditional expectation, so that we can reasonably guess that it moves around this value.¹

¹We remark that the log formulation of the intra-daily periodic component guarantees $\prod_{i=1}^I \phi_i = 1$, but not $\bar{\phi} = 1$. This is a minor issue, since $\bar{\phi}$ is a constant (once the periodic component is fixed) scaling factor for the daily component and is quite close to one in the applications considered.

Once the daily averages are computed, the ratio $x_{ti}^{(I)} = x_{ti}/\bar{x}_t$ can be used as a proxy for the whole intra-daily component $\phi_i \mu_{ti}$, since

$$x_{ti}^{(I)} = \frac{x_{ti}}{\bar{x}_t} \simeq \frac{\eta_t \phi_i \mu_{tj} \varepsilon_{ti}}{\eta_t \bar{\phi}} = \frac{\phi_i \mu_{tj} \varepsilon_{ti}}{\bar{\phi}}. \quad (10)$$

The bin average of the quantities into (10), namely $\bar{x}_{.i}^{(I)} = T^{-1} \sum_{t=1}^T x_{ti}^{(I)}$, represents a proxy for the intra-daily periodic component ϕ_i . In fact,

$$\bar{x}_{.i}^{(I)} = \frac{1}{T} \sum_{t=1}^T x_{tj}^{(I)} \simeq \frac{\phi_i}{\bar{\phi}} \frac{1}{T} \sum_{t=1}^T \mu_{tj} \varepsilon_{tj}. \quad (11)$$

By taking its expectation conditionally on the starting information, we have

$$E(\bar{x}_{.i}^{(I)} | \mathcal{F}_{0I}) \simeq \frac{\phi_i}{\bar{\phi}} \frac{1}{T} \sum_{t=1}^T E(\mu_{ti} | \mathcal{F}_{0I}) \simeq \frac{\phi_i}{\bar{\phi}}. \quad (12)$$

The last approximation can be motivated by considering that the average of the μ_{ti} 's for bin j converges, in some sense, to the unconditional average 1.

Finally, the residual quantity $x_{ti}^{(I)} / \bar{x}_{.i}^{(I)} = x_{ti} / \bar{x}_{.i}$ can be justified as a proxy for the intra-daily non-periodic component, since

$$\frac{x_{ti}^{(I)}}{\bar{x}_{.i}^{(I)}} \simeq \frac{\phi_i \mu_{tj} \varepsilon_{tj} / \bar{\phi}}{\phi_i / \bar{\phi}} = \mu_{ti} \varepsilon_{ti}. \quad (13)$$

3.2.2 CMEM and Component GARCH

Although the CMEM of Section 3.1 has some relationships with the component GARCH model suggested by Engle *et al.* (2007), for modeling intra-daily volatility, our proposal differs under a number of aspects. In particular, the main distinguishing feature lies in the evolution of the daily and intra-daily components: exploiting the scheme proposed, all parameters of the model can be estimated jointly, instead to recurring to a multi-step

procedure.

3.2.3 CMEM and Periodic GARCH

Our CMEM shares some features with the P-GARCH model (Bollerslev and Ghysels (1996)) as well. By grouping intra-daily components ϕ_i and μ_{ti} and referring to Equation (4) for the latter, the combined component can be written as

$$\phi_i \mu_{ti} = \alpha_{0i}^{(\mu)} + \beta_{1i}^{(\mu)} \mu_{ti-1} + \alpha_{1i}^{(\mu)} x_{ti-1}^{(\mu)}, \quad (14)$$

where

$$\alpha_{0i}^{(\mu)} = \alpha_0^{(\mu)} \phi_i \quad \alpha_{1i}^{(\mu)} = \alpha_1^{(\mu)} \phi_i \quad \beta_{1i}^{(\mu)} = \beta_1^{(\mu)} \phi_i. \quad (15)$$

In practice, the coefficients defined in (15) are periodic: their pattern is ruled by ϕ_i but each of them is rescaled by a (possibly) different value, as it would be in a P-GARCH-like formulation. We adopt a considerable simplification by imposing the same periodic pattern to all coefficients. In this respect, we are inspired by the results in Martens *et al.* (2002) where a relatively parsimonious formulation is adopted, based on an intra-daily periodic component scaling the dynamic (GARCH-like) component of the variance. The corresponding forecasts of the intra-daily volatility are only marginally worse than a more computationally expensive P-GARCH. Martens *et al.* (2002) provide also empirical evidence in favor of the exponential formulation of the periodic intra-daily component and support its representation in a Fourier form (even if in their application they consider only the first 4 harmonics). This notwithstanding, we substantially depart from their approach, since we add an explicit dynamic structure for the daily component, taking the intra-daily component to be a corresponding scale factor, and we estimate all CMEM parameters jointly.

3.3 Inference

Let us now illustrate how to obtain inferences on the model specified in Section 3.1. We group the main parameters of interest into the p -dimensional vector $\theta = (\theta^{(\eta)}; \phi; \theta^{(\mu)})$, where the three subvectors refer to the corresponding components of the model. Relative to these, the variance of the error term, σ^2 , represents a nuisance parameter.

Since the model is specified in a semiparametric way (see (1)), we focus our attention on the Generalized Method of Moments (GMM – Newey and McFadden (1994) and Wooldridge (1994)) as an estimation strategy not needing the adoption of a density function for the innovation term.

Rather than by GMM, MEMs are often estimated by QMLE by maximizing the log-likelihood of the specification based on a Gamma distribution assumption for the innovation term (see Engle and Gallo (2006)). The first order conditions for the conditional mean parameters are in fact the same for the two estimators. However, the portion of the Gamma log-likelihood due to the Gamma dispersion parameter is not defined or overflows numerically when, respectively, zeros or inliers² are present in the data.

By the same token, our GMM approach is robust to such features which are common in these datasets, especially when dealing with a higher number of intra-daily bins or illiquid assets.

3.3.1 Efficient GMM inference

Let

$$u_\tau = \frac{x_\tau}{\eta_t \phi_i \mu_\tau} - 1, \quad (16)$$

²Inliers are observations that are anomalous by being too small (in this context, too close to zero). Although relevant in general, in our empirical application this is not a problem.

where we simplified the notation by suppressing the reference to the dependency of u_τ on the parameters θ , on the information $\mathcal{F}_{\tau-1}$ and on the current value of the dependent variable x_τ . u_τ is a conditionally homoskedastic martingale difference, given that its conditional expectation is zero and its conditional variance is σ^2 . As a consequence, let us consider any $(M, 1)$ vector $\mathbf{G}_\tau$ depending deterministically on the information $\mathcal{F}_{\tau-1}$ and write $\mathbf{G}_\tau u_\tau \equiv \mathbf{g}_\tau$. We have

$$E(\mathbf{g}_\tau | \mathcal{F}_{\tau-1}) = \mathbf{0} \quad \forall \tau \Rightarrow E(\mathbf{g}_\tau) = \mathbf{0}, \quad (17)$$

by the law of iterated expectations; $\mathbf{g}_\tau$ is also a martingale difference.

Assuming that the absolute values of u_τ and $\mathbf{G}_\tau u_\tau$ have finite expectations, the uncorrelatedness of $\mathbf{G}_\tau$ and u_τ provides an instrument role to the former. $\mathbf{G}_\tau$ may depend on nuisance parameters collected into the vector ψ . In order for us to concentrate on estimating θ , we assume for the moment that ψ is a known constant, postponing any further discussion about its role and how to inference it to the end of this section and to Section 3.3.2.

If $M = p$, we have as many equations as the dimension of θ , thus leading to the moment criterion

$$\bar{\mathbf{g}} = \frac{1}{N} \sum_{\tau=1}^N \mathbf{g}_\tau = \mathbf{0}. \quad (18)$$

Under correct specification of the η_t , ϕ_i , and μ_{ti} equations and some regularity conditions, the GMM estimator $\hat{\theta}_N$, obtained solving (18) for θ , is consistent (Wooldridge (1994, th. 7.1)). Furthermore, under some additional regularity conditions, we have asymptotic normality of $\hat{\theta}_N$, with asymptotic covariance matrix (Wooldridge (1994, th. 7.2))

$$\text{Avar}(\hat{\theta}_N) = \frac{1}{N} (\mathbf{S}' \mathbf{V}^{-1} \mathbf{S})^{-1}, \quad (19)$$

where

$$\mathbf{S} = \lim_{N \rightarrow \infty} \frac{1}{N} \sum_{\tau=1}^N E(\nabla_{\theta'} \mathbf{g}_{\tau}) \quad (20)$$

$$\mathbf{V} = \lim_{N \rightarrow \infty} \frac{1}{N} V \left(\sum_{\tau=1}^N \mathbf{g}_{\tau} \right) = \lim_{N \rightarrow \infty} \left[\frac{1}{N} \sum_{\tau=1}^N E(\mathbf{g}_{\tau} \mathbf{g}_{\tau}') \right]. \quad (21)$$

The last expression for $\mathbf{V}$ comes from the fact that $\mathbf{g}_{\tau}$ is a martingale difference, since this is a sufficient condition for making these terms serially uncorrelated; moreover, the same condition leads to simplifications in the assumptions needed for the asymptotic normality, by virtue of the martingale CLT.

The martingale difference structure of u_{τ} gives also a simple formulation for the *efficient* choice of the instrument $\mathbf{G}_{\tau}$, associated with the 'smallest' asymptotic variance among the GMM estimators generated by $\bar{\mathbf{g}}$ functions structured as in (18). Such efficient choice is

$$\mathbf{G}_{\tau}^* = -E(\nabla_{\theta} u_{\tau} | \mathcal{F}_{\tau-1}) V(u_{\tau} | \mathcal{F}_{\tau-1})^{-1}. \quad (22)$$

Computing $E(\mathbf{g}_{\tau} \mathbf{g}_{\tau}')$ into (21) and $E(\nabla_{\theta'} \mathbf{g}_{\tau})$ into (20), we obtain

$$E(\mathbf{g}_{\tau} \mathbf{g}_{\tau}') = -E(\nabla_{\theta'} \mathbf{g}_{\tau}) = \sigma^2 E(\mathbf{G}_{\tau}^* \mathbf{G}_{\tau}^{*'}),$$

so that

$$\mathbf{V} = -\mathbf{S} = \sigma^2 \lim_{N \rightarrow \infty} \frac{1}{N} \sum_{\tau=1}^N E(\mathbf{G}_{\tau}^* \mathbf{G}_{\tau}^{*'})$$

and (19) specializes to

$$\text{Avar}(\hat{\theta}_N) = \frac{1}{N} (\mathbf{S}' \mathbf{V}^{-1} \mathbf{S})^{-1} = -\frac{1}{N} \mathbf{S}^{-1} = \frac{1}{N} \mathbf{V}^{-1}. \quad (23)$$

Considering the analytical structure of u_{τ} in the model (equation (16)), we have

$$\nabla_{\theta} u_{\tau} = -\mathbf{a}_{\tau}(u_{\tau} + 1),$$

where

$$\mathbf{a}_\tau = \eta_t^{-1} \nabla_{\boldsymbol{\theta}} \eta_t + \mu_\tau^{-1} \nabla_{\boldsymbol{\theta}} \mu_\tau + \phi_i^{-1} \nabla_{\boldsymbol{\theta}} \phi_i \quad (24)$$

so that (22) becomes

$$\mathbf{G}_\tau^* = \mathbf{a}_\tau \sigma^{-2}.$$

Replacing it into $\mathbf{g}_\tau = \mathbf{G}_\tau u_\tau$ and this, in turn, into (18), we obtain that the GMM estimator of $\boldsymbol{\theta}$ in the CMEM solves the MM equation

$$\frac{1}{N} \sum_{\tau=1}^N \mathbf{a}_\tau u_\tau = \mathbf{0}, \quad (25)$$

which does not depend on the nuisance parameter σ^2 and, therefore, inference relative to the main parameter $\boldsymbol{\theta}$ does not depend on the estimation of σ^2 .

The asymptotic variance matrix of $\hat{\boldsymbol{\theta}}_N$ is

$$\text{Avar}(\hat{\boldsymbol{\theta}}_N) = \frac{\sigma^2}{N} \left[\lim_{N \rightarrow \infty} \frac{1}{N} \sum_{\tau=1}^N E(\mathbf{a}_\tau \mathbf{a}_\tau') \right]^{-1} \quad (26)$$

that can be consistently estimated by

$$\widehat{\text{Avar}}(\hat{\boldsymbol{\theta}}_N) = \hat{\sigma}_N^2 \left[\sum_{\tau=1}^N \mathbf{a}_\tau \mathbf{a}_\tau' \right]^{-1} \quad (27)$$

where $\hat{\sigma}_N^2$ is a consistent estimator of σ^2 (Section 3.3.2) and $\mathbf{a}_\tau$ is here evaluated at $\hat{\boldsymbol{\theta}}_N$.

3.3.2 Inference on σ^2

A straightforward estimator for the second moment σ^2 of u_τ (cf. (16)), which is not compromised by zeros in the data, is

$$\hat{\sigma}_N^2 = \frac{1}{N} \sum_{\tau=1}^N \hat{u}_\tau^2 \quad (28)$$

where $\hat{u}_\tau$'s are computed from current values of $\hat{\theta}_N$.

3.4 Empirical Application: In Sample Volume Analysis

The empirical application focuses on the same SPY, DIA and QQQQ tickers previously examined, dividing each trading day in equally spaced bins of 15 minutes. We make the model more realistic by inserting some additional features which would have encumbered the notation presented before. We allow for a specific behavior in the first bin of the day to accommodate the possible accumulation of news during closing time: a higher level of turnover at the beginning of the day is not necessarily linked to market activity in the last bin of the previous day (see Figure 2). The insertion of a dummy variable for the first bin into the dynamic intra-daily component (in Equation (4) with a coefficient $\nu_1^{(\mu)}$) marks the break in trading.³ In order to allow for possible differences in the dynamics related to the sign of returns we also add asymmetric effects to the right hand side of the daily and intra-daily dynamic components (Equations (2) and (4), respectively)

$$\gamma_1^{(\eta)} x_{t-1}^{-(\eta)} \quad \text{and} \quad \gamma_1^{(\mu)} x_{ti-1}^{-(\mu)},$$

where

$$x_t^{-(\eta)} = x_t^{(\eta)} \mathbf{I}(r_t < 0) \quad \text{and} \quad x_{ti}^{-(\mu)} = x_{ti}^{(\mu)} \mathbf{I}(r_{ti} < 0)$$

denote the asymmetric versions of the deflated daily and intra-daily volumes, respectively, as defined in Section 3.1. r_t is the total return in day t and r_{ti} is the return in bin i of day t (we assume that the return distribution is symmetric around zero). The inclusion of asymmetric effects can be motivated by invoking the well documented greater impact of bad news on volatility (*leverage effect* – Nelson (1991), Glosten *et al.* (1993), Rabemananjara and Zakoïan (1993)) and the recognized existence of a common latent component (commonly interpreted as *information flow*) behind both volatilities and volumes (see Andersen

³We thank Torben Andersen for pointing this out to us.

(1996), Hautsch (2008)). Because of this, it is reasonable to expect a positive impact of negative returns on volumes as well. Finally, we allow for a second lag in the intra-daily dynamic component.

Summarizing, by mixing the previous ingredients we consider the following CMEMs:

base: CMEM with a dummy at bin 1, lag-1 dependence and no asymmetric effects;

asym: base CMEM with lag-1 asymmetric effects (daily and intra-daily);

intra2: base CMEM with a second lag in the intra-daily dynamic component;

asym-intra2: intra2 CMEM with lag-1 asymmetric effects (daily and intra-daily).

The parameter estimates of the daily and intra-daily dynamic components are reported in Table 2. $\alpha_0^{(\mu)}$ lacks standard errors because it is estimated via expectation targeting by imposing $E(\mu_\tau) = 1$ (as detailed in section 3.1); the remaining parameters are jointly estimated via GMM as detailed in Section 3.3.

The parameter estimates are remarkably similar across assets, suggesting some common behavior in the volume dynamics. In the present context $\alpha_1^{(\eta)}$ is much larger than the customary values encountered in typical GARCH(1,1) estimated on daily returns. As expected, the coefficient of the daily asymmetric effect ($\gamma^{(\eta)}$) is positive and always significant; moreover, its inclusion tends in general to increase $\beta^{(\eta)}$ slightly. The level of persistence of the daily component ($\alpha_1^{(\eta)} + \gamma^{(\eta)}/2 + \beta^{(\eta)}$) is in general around 0.95.

The level of persistence in the intra-daily component (column $\text{pers}(\mu)$), evaluated as the largest eigenvalue of the companion matrix built from $\alpha_l^{(\eta)} + \gamma_l^{(\eta)}/2 + \beta_l^{(\eta)}$ estimates, $l = 1, 2$) is relatively high and increases remarkably when $\alpha_2^{(\mu)}$ is present. This last parameter is negative and with a relatively large magnitude, but it is such that the Nelson and Cao (1992) non-negativity condition for the corresponding component is satisfied in all cases. The coefficients of the intra-daily asymmetric effects are positive and significant, even if smaller in magnitude than the corresponding daily values.

Table 3 displays some diagnostics on the estimated residuals. Summarizing, the CMEM makes a relatively good job in modeling turnover, in particular in the presence of a second lag in the intra-daily dynamic component. For such models, Ljung-Box statistics at lag 1 reveal that only the DIA ticker shows some marginally significant autocorrelation in residuals. More problems are rather at lag 1-day, for which only DIA gives insignificant statistics. To be noted that the very large number of observations makes the statistic sensitive to even small departures from the uncorrelatedness assumption. This appears to be confirmed by Figure 5 (relative to SPY residuals) for which single autocorrelations are substantially within the confidence bands whereas the global statistic is significant. On the other hand, squared residuals are by far less problematic for all formulations taken into account. The table and the figure reveal that the contribution of the lag-2 autoregressive term in improving the diagnostics is more relevant than the one of the asymmetric terms.

Finally, in order to provide some idea about the impact of different intra-daily frequencies, we estimated the same formulations on the tickers just considered but with bins of 30 minutes instead of 15 (Table 4). Focusing on the asym-intra2 version, the intra-daily persistence tends to diminish for larger bins, in particular as consequence of smaller $\beta^{(\mu)}$ values, whereas intra-daily asymmetric effects tend to increase slightly. Also the coefficients of the daily component tend to adjust as an effect of the different sampling frequency. The persistence is about the same, but this comes from larger $\beta^{(\eta)}$'s and smaller $\alpha_1^{(\eta)}$'s. The doubling of the $\alpha_0^{(\eta)}$ coefficients is just due to the fact that halving the number of the intra-daily bins (from 26 to 13) the average level tends to double.




Figure 5: SPY Turnover Data: Residuals. January 2002 to December 2006.

Ticker	Specification	$\alpha_0^{(\eta)}$	$\alpha_1^{(\eta)}$	$\gamma^{(\eta)}$	$\beta^{(\eta)}$	$\alpha_0^{(\mu)}$	$\nu_1^{(\mu)}$	$\alpha_1^{(\mu)}$	$\alpha_2^{(\mu)}$	$\gamma^{(\mu)}$	$\beta^{(\mu)}$	σ	pers(μ)
SPY	base	1.811 (0.369)	0.474 (0.037)		0.483 (0.040)	0.230	0.914 (0.207)	0.340 (0.007)			0.395 (0.013)	0.622	0.735
	asym	1.665 (0.348)	0.431 (0.037)	0.041 (0.010)	0.511 (0.039)	0.227	0.968 (0.213)	0.317 (0.008)		0.036 (0.006)	0.400 (0.013)	0.622	0.735
	intra2	1.925 (0.458)	0.504 (0.052)		0.452 (0.057)	0.034	1.056 (0.172)	0.354 (0.008)	-0.235 (0.010)		0.807 (0.012)	0.622	0.900
	asym-intra2	1.795 (0.432)	0.461 (0.050)	0.041 (0.012)	0.480 (0.054)	0.034	1.084 (0.170)	0.338 (0.008)	-0.233 (0.010)	0.024 (0.004)	0.808 (0.012)	0.621	0.899
DIA	base	1.951 (0.420)	0.419 (0.035)		0.536 (0.039)	0.199	0.897 (0.215)	0.283 (0.008)			0.484 (0.013)	0.834	0.767
	asym	1.881 (0.387)	0.361 (0.035)	0.06 (0.012)	0.567 (0.038)	0.200	0.934 (0.222)	0.265 (0.008)		0.03 (0.008)	0.484 (0.013)	0.833	0.751
	intra2	1.766 (0.479)	0.452 (0.045)		0.511 (0.049)	0.041	1.237 (0.223)	0.309 (0.009)	-0.188 (0.011)		0.791 (0.013)	0.833	0.888
	asym-intra2	1.749 (0.449)	0.397 (0.044)	0.057 (0.014)	0.538 (0.047)	0.041	1.235 (0.222)	0.295 (0.009)	-0.187 (0.010)	0.022 (0.005)	0.793 (0.013)	0.832	0.889
QQQQ	base	1.537 (0.369)	0.466 (0.036)		0.507 (0.038)	0.218	1.134 (0.220)	0.359 (0.007)			0.380 (0.012)	0.538	0.739
	asym	1.466 (0.350)	0.423 (0.036)	0.045 (0.008)	0.530 (0.037)	0.222	1.145 (0.224)	0.351 (0.008)		0.012 (0.006)	0.376 (0.012)	0.538	0.733
	intra2	1.815 (0.460)	0.485 (0.052)		0.484 (0.055)	0.026	1.085 (0.157)	0.372 (0.007)	-0.258 (0.01)		0.818 (0.011)	0.536	0.905
	asym-intra2	1.746 (0.440)	0.448 (0.051)	0.041 (0.01)	0.503 (0.054)	0.027	1.119 (0.162)	0.365 (0.008)	-0.256 (0.01)	0.011 (0.003)	0.815 (0.011)	0.536	0.902

Table 2: Parameter estimates for CMEM on turnover data. Sample period 2002 – 2006; intra-daily bins taken at 15 minutes. Standard errors are reported in parentheses. pers(μ) indicates estimated persistence of the dynamic intra-daily component.

Ticker	Model	Residuals $\hat{\epsilon}_{ti}$			Squared Residuals $\hat{\epsilon}_{ti}^2$				
		$\hat{\rho}_1$	$\hat{\rho}_1$ day	Q_1	Q_1 day	$\hat{\rho}_1$	$\hat{\rho}_1$ day	Q_1	Q_1 day
SPY	base	0.012	0.014	4.847 0.028	179.291 0.000	0.000	0.006	0.004 0.948	36.406 0.084
	asym	0.013	0.014	5.584 0.018	177.625 0.000	0.000	0.006	0.000 0.988	39.038 0.048
	intra2	-0.007	0.013	1.582 0.208	47.936 0.005	-0.005	0.005	0.937 0.333	37.450 0.068
	asym-intra2	-0.006	0.013	1.354 0.245	49.87 0.003	-0.005	0.005	0.970 0.325	39.536 0.043
DIA	base	0.011	0.004	3.604 0.058	147.463 0.000	-0.003	0.006	0.251 0.616	17.850 0.881
	asym	0.010	0.004	3.331 0.068	153.42 0.000	-0.003	0.005	0.247 0.886	17.725 0.619
	intra2	-0.013	0.005	5.328 0.021	36.335 0.086	-0.008	0.007	1.869 0.172	11.787 0.992
	asym-intra2	-0.012	0.005	5.004 0.025	36.103 0.090	-0.008	0.008	1.887 0.170	12.233 0.990
QQQQ	base	0.019	0.002	11.647 0.001	239.935 0.000	0.011	-0.003	4.062 0.044	52.965 0.001
	asym	0.020	0.002	12.411 0.000	250.289 0.000	0.012	-0.003	4.997 0.025	55.825 0.001
	intra2	-0.001	0.000	0.04 0.841	61.145 0.000	0.006	-0.004	1.262 0.261	34.250 0.129
	asym-intra2	-0.001	0.000	0.046 0.830	59.348 0.000	0.005	-0.004	0.864 0.353	32.806 0.168

Ticker	Specification	$\alpha_0^{(\eta)}$	$\alpha_1^{(\eta)}$	$\gamma^{(\eta)}$	$\beta^{(\eta)}$	$\alpha_0^{(\mu)}$	$\nu_1^{(\mu)}$	$\alpha_1^{(\mu)}$	$\alpha_2^{(\mu)}$	$\gamma^{(\mu)}$	$\beta^{(\mu)}$	σ	$\text{pers}(\mu)$
SPY	asym-intra2	3.011 (0.786)	0.348 (0.040)	0.045 (0.011)	0.597 (0.044)	0.055	0.999 (0.161)	0.299 (0.011)	-0.172 (0.015)	0.037 (0.005)	0.723 (0.024)	0.532	0.813
DIA	asym-intra2	2.968 (0.804)	0.314 (0.037)	0.062 (0.013)	0.625 (0.040)	0.065	1.125 (0.2)	0.268 (0.011)	-0.138 (0.015)	0.043 (0.007)	0.697 (0.025)	0.678	0.795
QQQQ	asym-intra2	2.945 (0.816)	0.336 (0.041)	0.044 (0.009)	0.619 (0.044)	0.046	1.081 (0.167)	0.323 (0.011)	-0.184 (0.016)	0.021 (0.005)	0.722 (0.025)	0.47	0.824

Table 4: Parameter estimates for CMEM on turnover data. Sample period 2002 – 2006; intra-daily bins taken at 30 minutes. Standard errors are reported in parentheses. $\text{pers}(\mu)$ indicates estimated persistence of the dynamic intra-daily component.

4 Intra-daily Volume Forecasting for VWAP Trading

4.1 VWAP Trading

The concept of daily Volume Weighted Average Price (VWAP) was introduced by Berkowitz *et al.* (1988)) as an average of intra-daily transaction prices weighted by the corresponding traded volume relative to the total volume traded during the day (*full VWAP* in Madhavan (2002)). In the original paper, the difference between the price of a trade and the recorded VWAP was used to measure the market impact of that trade. As such, VWAP is used to evaluate execution performance, given that it is a very transparent measure, easily calculated at the end of the day with tick-by-tick data.

VWAP trading is thus defined as a procedure for splitting a certain number of shares into smaller size orders during the day, which will be executed at different prices with the net result of an average price that is close to the VWAP⁴. In what follows, the trade to be executed is treated as exogenously determined (cf. Bertsimas and Lo (1998), Almgren and Chriss (2000), Engle and Ferstenberg (2007)). In order to implement the replicating strategy, we assume that we are price takers and no effort will be put in predicting prices while we concentrate on the fact that accurate intra-daily volume proportion forecasting is the base for VWAP replication.

Let the VWAP for day t be defined as

$$\text{VWAP}_t = \frac{\sum_{j=1}^{J_t} v_t(j) p_t(j)}{\sum_{j=1}^{J_t} v_t(j)}.$$

where $p_t(j)$ and $v_t(j)$ denote, respectively, the price and volume of the j -th transaction of day t and J_t is the total number of trades of day t . For a given partition of the trading day

⁴Whether the VWAP benchmark is proposed on an agency base or on a guaranteed base (in exchange for a fee) is a technical aspect which does not have any bearings in what we discuss.

into I bins, it is possible to express the numerator of the VWAP as

$$\begin{aligned}\sum_{j=1}^{J_t} v_t(j) p_t(j) &= \sum_{i=1}^I \left(\sum_{j \in \mathcal{J}_i} v_t(j) \right) \bar{p}_{ti} \\ &= \sum_{i=1}^I x_{ti} \bar{p}_{ti},\end{aligned}$$

where $\bar{p}_{ti}$ is the VWAP of the i -th bin and $\mathcal{J}_i$ denotes the set of indices of the trades belonging to the i -th bin. Hence,

$$\text{VWAP}_t = \frac{\sum_{i=1}^I x_{ti} \bar{p}_{ti}}{\sum_{i=1}^I x_{ti}} = \sum_{i=1}^I w_{ti} \bar{p}_{ti} = \mathbf{w}'_t \bar{\mathbf{p}}_t$$

where w_{ti} is the intra-daily proportion of volumes traded in bin i on day t , that is $w_{ti} = x_{ti} / \sum_{i=1}^I x_{ti}$. Let $\mathbf{y} = (y_1, \dots, y_I)$, an order slicing strategy over day t with the same bin intervals. We can define the Average Execution Price as the quantity

$$\text{AEP}_t = \sum_{i=1}^I y_i \bar{p}_{ti} = \mathbf{y}' \bar{\mathbf{p}}_t.$$

The choice variable being the vector $\mathbf{y}$, we can solve the problem of minimizing the distance between the two outcomes in a mean square error sense, namely

$$\min_{\mathbf{y}} \delta_t = (\mathbf{w}'_t \bar{\mathbf{p}}_t - \mathbf{y}' \bar{\mathbf{p}}_t)^2,$$

where, solving the minimization problem leads to the first order conditions

$$\frac{d\delta_t}{d\mathbf{y}} = \mathbf{0} \Rightarrow -2\bar{\mathbf{p}}_t(\mathbf{w}_t - \mathbf{y})' \bar{\mathbf{p}}_t = \mathbf{0}, \quad (29)$$

which has a meaningful solution for $\mathbf{y} = \mathbf{w}_t$, that is, when the order slicing sequence for each sub period in the day reproduces exactly the overall relative volume for that sub period.

The implication of Equation (29) is that the VWAP replication problem can be cast as an intra-daily volume proportion forecasting problem: the better we can predict the intra-daily volumes proportions, the better we can track VWAP.

4.2 VWAP Replication Strategies

Following Białkowski *et al.* (2008), we consider two types of VWAP replication strategies: Static and Dynamic. The Static VWAP replication strategy assumes that the order slicing is set before the market opening and it is not revised during the trading day. In the Dynamic VWAP replication strategy scenario, on the other hand, order slicing is revised at each new sub period as new intra-daily volumes are observed.

Let $\hat{x}_{t i|t-1}$ be shorthand notation for the prediction of $x_{t i}$ conditionally on the previous day full information set $\mathcal{F}_{t-1 I}$. The *Static VWAP* replication strategy is implemented using slices with weights given by

$$\hat{w}_{t i|t-1} = \frac{\hat{x}_{t i|t-1}}{\sum_{i=1}^I \hat{x}_{t i|t-1}} \quad i = 1, \dots, I,$$

that is, the proportion of volumes for bin i is given by predicted volume in bin i divided by the sum of the volume predictions.

Let $\hat{x}_{t i|i-1}$ be shorthand notation to denote the prediction of $x_{t i}$ conditionally on $\mathcal{F}_{t i-1}$. The *Dynamic VWAP* replication strategy is implemented using slices with weights given by

$$\hat{w}_{t i|i-1} = \begin{cases} \frac{\hat{x}_{t i|i-1}}{\sum_{j=1}^I \hat{x}_{t j|i-1}} \left(1 - \sum_{j=1}^{i-1} \hat{w}_{t j|j-1}\right) & i = 1, \dots, I-1 \\ \left(1 - \sum_{i=1}^{I-1} \hat{w}_{t i|i-1}\right) & i = I \end{cases}$$

that is, for each intra-daily bin from 1 to $I-1$ the predicted proportion is given by the proportion of 1-step ahead volumes with respect to the sum of the remaining predicted volumes multiplied by the slice proportion left to be traded. On the last period of the




Figure 6: Slicing Loss function for $I = 3$ and $(w_1, w_2, w_3) = (0.3, 0.3, 0.4)$.

day I , the predicted proportion is equal to the remaining part of the slice that needs to be traded.

4.3 Forecast Evaluation

We evaluate out-of-sample performance from different perspectives: intra-daily volumes, intra-daily volume proportions and daily VWAP prediction.

A natural way to assess volume predictive ability is to consider the mean square prediction error of the volume forecasts, defined as

$$\text{MSE}^{\text{vol}} = \sum_{t=1}^T \sum_{i=1}^I (x_{ti} - \hat{x}_{ti|\cdot})^2,$$

where $\hat{x}_{ti|\cdot}$ denotes the volume from some VWAP replication and volume forecasting strategy. Although such a metric provides insights as to which model gives a more realistic description of volume dynamics, it does not necessarily provide useful information as to

the performance of the models for VWAP trading.

The evaluation of the accuracy of intra-daily volume proportion forecasts calls for a suitable loss function, given the bounded interval within which they are constrained. As no well established criteria to measure proportion predictions ability are present in the literature, we propose the following *Slicing* loss function:

$$L^{\text{slicing}} = - \sum_{t=1}^T \sum_{i=1}^I w_{ti} \log \hat{w}_{ti}, \quad (30)$$

which we motivate from both an “operational” as well as an information theoretic perspective. In the spirit of Christoffersen (1998) in the context of Value at Risk, we ask ourselves which properties proportion forecast ought to have under correct specification. Assume that a broker is interested in trading n shares of the asset⁵ each day. If the intra-daily volume proportion predictions are correct, the observed intra-daily volumes nw_{ti} , $i = 1, \dots, I$ behave like a sample from a multinomial distribution with parameters $\hat{w}_{ti}$, $i = 1, \dots, I$, and n ; that is

$$(nw_{t1}, \dots, nw_{tI}) \sim \text{Mult}(\hat{w}_{t1}, \dots, \hat{w}_{tI}, n).$$

This suggests that an appropriate loss function for the evaluation of such forecasts is the negative of the multinomial predictive log-likelihood

$$L^{\text{Mult}} = - \sum_{t=1}^T \left(\log \frac{n!}{(nw_{t1})! \dots (nw_{tI})!} + \sum_{i=1}^I nw_{ti} \log \hat{w}_{ti} \right).$$

An alternative evaluation strategy consists of computing the distance between the actual and predicted intra-daily volume proportions as the discrepancy between two discrete

⁵We are implicitly assuming for simplicity's sake, that the actual intra-daily proportions nw_{ti} are all integer.

distributions. Using the Kullback–Leibler measure we get

$$L^{\text{KL}} = \sum_{t=1}^T \sum_{i=1}^I (w_{ti} \log w_{ti} - w_{ti} \log \widehat{w}_{ti}).$$

Interestingly, both the Multinomial and Kullback–Leibler losses provide equivalent rankings among competing forecasting methods in that the comparison is driven by the common term $-\sum_{t=1}^T \sum_{i=1}^I w_{ti} \log \widehat{w}_{ti}$. Figure 6 shows a picture of the Slicing loss function in the case of 3 intra-daily bins when the actual proportions $\mathbf{w}_t$ are $(0.3, 0.3, 0.4)$. The Slicing loss function is defined over the $I - 1$ dimensional simplex described by $\sum_{i=1}^I \widehat{w}_{ti} = 1$, has a minimum in correspondence to the true values, the value of the loss function goes to infinity on the boundaries of the simplex (when the actual proportions are in the interior of the simplex) and is asymmetric.

Finally, we also consider *VWAP tracking errors* MSE as in Białkowski *et al.* (2008) defined as

$$\text{MSE}^{\text{VWAP}} = \sum_{t=1}^T \left(\frac{\text{VWAP}_t - \widehat{\text{VWAP}}_t}{\text{VWAP}_t} 100 \right)^2,$$

where VWAP_t is the VWAP of day t and $\widehat{\text{VWAP}}_t$ is the realized average execution price obtained using some VWAP replication strategy and volume forecasting method. Both VWAP_t and $\widehat{\text{VWAP}}_t$ are computed using the last recorded price of the i -th bin as a proxy for the average price of the same interval. The VWAP tracking error for day t can be seen as an average of slicing errors within each bin weighed by the relative deviation of the price associated to that bin with respect to the VWAP:

$$\text{MSE}^{\text{VWAP}} = \sum_{t=1}^T \left(\sum_{i=1}^I (w_{ti} - \widehat{w}_{ti}) \frac{\bar{p}_{ti}}{\text{VWAP}_t} \right)^2 100^2$$

Note that the deviations of the prices from the daily VWAP add an extra source of noise which can spoil the correct ranking of slice forecasts. In light of this, we recommend




Figure 7: SPY: sequences of multi step ahead volume predictions. asym-intra2 specification starting from different intra-daily bins (0, 6, 12 and 19) until the end of the day. January 31st, 2005.

evaluating the precision of the forecasts by means of the Slicing loss function.

4.4 Empirical Application: Out-of-Sample VWAP Prediction

Our empirical application consists of a volume, volume proportion and VWAP tracking exercise of the tickers SPY, DIA and QQQQ between January 2005 and December 2006 (502 days, 6526 observations). Model mis-specification and parameter instability is always a possible concern in forecasting: to this end, we adopt parameter estimates for our CMEM specifications which are recursively updated each week using data from 2002, producing turnover predictions under both Static and Dynamic VWAP replication strategies. In order to assess the usefulness of the proposed approach we use (periodic) Rolling Means (RM) as a simple benchmark, that is the predicted volume for the i -th bin is obtained as the mean over the last 40 days at the same bin. The Rolling Means are implemented within the Static VWAP replication approach.

	SPY			DIA			QQQ		
	volume	slicing	vwap	volume	slicing	vwap	volume	slicing	vwap
RM	43.179	3.203	1.116	46.167	2.5163	1.334	53.296	3.184	1.799
base	41.623***	3.198***	1.114	45.124***	3.198***	1.350	51.740***	3.180***	1.785
asym	41.537***	3.198***	1.114	45.039***	3.198***	1.351	51.571***	3.180***	1.783
intra2	41.623***	3.197***	1.101	45.124***	3.197***	1.335	51.740***	3.179***	1.784
asym-intra2	41.537***	3.197***	1.097	45.039***	3.197***	1.335	51.571***	3.179***	1.785
					Dynamic				
base	36.958***	3.192***	1.101***	41.843***	3.192***	1.222***	46.514***	3.175***	1.772**
asym	36.939***	3.192***	1.099***	41.818***	3.192***	1.217***	46.459***	3.175***	1.767**
intra2	36.924***	3.189***	1.074***	41.710***	3.189***	1.222***	46.193***	3.171***	1.739***
asym-intra2	36.912***	3.189***	1.072***	41.696***	3.189***	1.219***	46.148***	3.171***	1.737***

Table 5: Out-of-Sample Volume, Slicing and VWAP tracking forecasting results. For each ticker, specification and VWAP replication strategy the table reports the values of the Volume, Slicing and VWAP tracking error loss functions. Asterisks denote the significance (* 1%, ** 5% and *** 10%) of a Diebold-Mariano test of equal predictive ability with respect to RM using the corresponding loss functions.

Figure 7 displays sequences of volume predictions for the SPY ticker on January 31, 2005. The forecast sequences are produced using the asym-intra2 specification starting from different intra-daily bins until the end of the day. We use this picture to offer some informal remarks on the potential differences in slicing decisions generated by the static and dynamic VWAP replication schemes. The periodic intra-daily pattern appears to dominate the intra-daily evolution of turnover forecasts but intra-daily bursts of activity can alter the intra-daily volume profile. This suggests that it is not straightforward to see large differences in the order slicing in the static VWAP replication case between our models and the benchmark. On the other hand, in the dynamic VWAP replication case, activity bursts can lead to quite improved order slicing decisions.

Table 5 reports the volume MSE, slicing loss and VWAP tracking MSE together with asterisks denoting the significance of a Diebold-Mariano test of equal predictive ability with respect to RM using the corresponding loss functions. In terms of volume and volume proportion predictions, the CMEM dynamic VWAP replication performs best and significantly outperforms the benchmark, followed by the CMEM static VWAP replication which generally outperforms the benchmark as well. The ranking of the CMEM specifications reflects the in-sample estimation results with models with richer intra-daily dynamics and asymmetric terms performing the best. Results also point out that it is the inclusion of an extra intra-daily lag rather than asymmetric effects which delivers the most of the out-of-sample gains. Recalling the words of caution about its limits, the VWAP tracking MSE evidence is substantially in line with the volume MSE and slicing loss results. In the static case, our models do systematically better than the benchmark in the QQQQ and SPY case, although significance of outperformance is lost. In the dynamic case our CMEMs are always able to significantly beat the benchmark, with the richest specifications obtaining the best results.

5 Conclusions

In this paper we suggest a dynamic model with different components capturing the behavior of traded volumes (relative to outstanding shares) viewed from daily and (periodic and non-periodic) intra-daily time perspectives. The parameters of this Component Multiplicative Error Model can be estimated in one step by the Generalized Method of Moments. The application to three major ETFs shows that both the static and the dynamic VWAP replication strategies generally outperform a commonly used naïve method of rolling means for intra-daily volumes in an out-of sample forecasting exercise.⁶

While focused here on specific assets, the analysis can be extended to uncover what other features to the model may be suggested by other classes of assets. For individual companies, for example, we may want to see if some stock specific characteristics (e.g. market capitalization, debt-to-equity ratio or percentage of holdings by institutional investors) have a bearing on the characteristics of the estimated dynamics.

The CMEM can be used in other contexts in which intra-daily bins are informative of some periodic features (e.g. volatility, number of trades, average durations) together with overall dynamics having components at different frequencies. The periodic component can be more parsimoniously specified by recurring to some shrinkage estimation as in Brownlees and Gallo (2008). Multivariate extensions are also possible (following Cipollini *et al.* (2009) by retrieving a richer price-volume dynamics in order to establish a relationship that can be related to the flow of information at different frequencies, separating it from (possibly common) periodic components.

⁶While statistically significant, we leave an open issue as to whether the improvement of our procedure could be documented in terms of economic savings it entails in pursuing the correct volume forecast. This goes beyond the scope of this paper as a simulation should be based on a joint model of price and volume determination with market microstructure foundations. On the other hand, evaluating the procedure on an actual exchange would be limited by the lack of access to an algorithmic trading system and in any case by proprietary limitations in diffusing the results.

References

- Almgren, R. and Chriss, N. (2000). Optimal execution of portfolio transactions. *Journal of Risk*, **3**, 5–39.
- Andersen, T. G. (1996). Return volatility and trading volume: An information flow interpretation of stochastic volatility. *The Journal of Finance*, **51**, 169–204.
- Berkowitz, S. A., Logue, D. E., and Noser, E. A. J. (1988). The total cost of transactions on the NYSE. *The Journal of Finance*, **43**(1), 97–112.
- Bertsimas, D. and Lo, A. W. (1998). Optimal control of execution costs. *Journal of Financial Markets*, **1**, 1–50.
- Białkowski, J., Darolles, S., and Le Fol, G. (2008). Improving vwap strategies: A dynamic volume approach. *Journal of Banking and Finance*, **32**, 1709–1722.
- Bikker, J., Spierdijk, L., Hoevenaars, R., and van der Sluis, P. J. (2008). Forecasting market impact costs and identifying expensive trades. *Journal of Forecasting*, **27**, 21–39.
- Bollerslev, T. and Ghysels, E. (1996). Periodic autoregressive conditional heteroskedasticity. *Journal of Business and Economic Statistics*, **14**, 139–151.
- Brownlees, C. T. and Gallo, G. M. (2006). Financial econometric analysis at ultra-high frequency: Data handling concerns. *Computational Statistics and Data Analysis*, **51**, 2232–2245.
- Brownlees, C. T. and Gallo, G. M. (2008). Shrinkage estimation of semiparametric multiplicative error models. Technical report, Dipartimento di Statistica “G.Parenti” – Florence.
- Chordia, T., Roll, R., and Subrahmanyam, A. (2008). Why has trading volume increased? Technical report, UCLA.

- Christoffersen, P. F. (1998). Evaluation of interval forecasts. *International Economic Review*, **39**, 841–862.
- Cipollini, F., Engle, R. F., and Gallo, G. M. (2009). Semiparametric vector MEM. Econometrics Working Papers Archive wp200903, Universita' degli Studi di Firenze, Dipartimento di Statistica "G. Parenti".
- Engle, R. F. (2002). New frontiers for ARCH models. *Journal of Applied Econometrics*, **17**, 425–446.
- Engle, R. F. and Ferstenberg, R. (2007). Execution risk. *Journal of Portfolio Management*, **33**, 34–45.
- Engle, R. F. and Gallo, G. M. (2006). A multiple indicators model for volatility using intra-daily data. *Journal of Econometrics*, **131**, 3–27.
- Engle, R. F., Ferstenberg, R., and Russell, J. (2006). Measuring and modeling execution cost and risk. Technical Report FIN-06-044, Stern - NYU.
- Engle, R. F., Sokalska, M. E., and Chanda, A. (2007). Forecasting intraday volatility in the US equity market. Multiplicative component GARCH. Technical report, North American Winter Meetings of the Econometric Society.
- Gallant, A. R. (1981). On the bias in flexible functional forms and an essentially unbiased form: The Fourier flexible form. *Journal of Econometrics*, **15**, 211–245.
- Gallant, A. R. (1984). The Fourier flexible form. *American Journal of Agricultural Economics*, **66**, 204–208.
- Glosten, L. R., Jagannathan, R., and Runkle, D. E. (1993). On the relation between the expected value and the volatility of the nominal excess return on stocks. *The Journal of Finance*, **48**(5), 1779–1801.
- Gouriéroux, C. S., Jasiak, J., and Le Fol, G. (1999). Intra-day market activity. *Journal of Financial Markets*, **2**(3), 193–226.

- Härdle, W. K., Hautsch, N., and Mihoci, A. (2009). Modelling and forecasting liquidity supply using semiparametric factor dynamics. Technical report, 2009-044, CRC 649, Humboldt Universität.
- Hautsch, N. (2008). Capturing common components in high-frequency financial time series: A multivariate stochastic multiplicative error model. *Journal of Economic Dynamics and Control*, **32**(12), 3978 – 4015.
- Madhavan, A. (2002). *VWAP Strategies*, pages 32–38. Investment Guides Series. Institutional Investor Inc.
- Martens, M., Chang, Y.-C., and Taylor, S. J. (2002). A comparison of seasonal adjustment methods when forecasting intraday volatility. *Journal of Financial Research*, **25**(2), 283–299.
- Nelson, D. B. (1991). Conditional heteroskedasticity in asset returns: A new approach. *Econometrica*, **59**, 347–370.
- Nelson, D. B. and Cao, C. Q. (1992). Inequality constraints in the univariate GARCH model. *Journal of Business and Economic Statistics*, **10**, 229–235.
- Newey, W. K. and McFadden, D. (1994). Large sample estimation and hypothesis testing. In R. F. Engle and D. McFadden, editors, *Handbook of Econometrics*, volume 4, chapter 36, pages 2111–2245. Elsevier.
- Rabemananjara, R. and Zakoïan, J. M. (1993). Threshold ARCH models and asymmetries in volatility. *Journal of Applied Econometrics*, **8**(1), 31–49.
- Taylor, N. (2002). The economic and statistical significance of spread forecasts: Evidence from the London Stock Exchange. *Journal of Banking and Finance*, **26**, 795–818.
- White, H. (2006). Approximate Nonlinear Forecasting Methods. In G. Elliott, C. W. J. Granger, and A. Timmermann, editors, *Handbook of Economic Forecasting*, volume 1, chapter 9, pages 459–512. Elsevier.

Wooldridge, J. M. (1994). Estimation and inference for dependent processes. In R. F. Engle and D. McFadden, editors, *Handbook of Econometrics*, volume 4, chapter 45, pages 2639–2738. Elsevier.